

Yours Truly

Krantz Template



*To my dog
and my cat.*



Contents

Foreword	ix
Preface	xi
List of Figures	xiii
List of Tables	xv
Contributors	xvii
Symbols	xxi
I This is What a Part Would Look Like	1
1 Basic Concepts	3
1.1 Introduction	3
1.1.1 A component part	4
1.2 Record Linkage Model	8
1.2.1 A component part	8
1.2.1.1 H3 A component part	8
1.3 Glossary	12
Bibliography	13
Index	15



Foreword

I am delighted to introduce the first book on Multimedia Data Mining. When I came to know about this book project undertaken by two of the most active young researchers in the field, I was pleased that this book is coming in early stage of a field that will need it more than most fields do. In most emerging research fields, a book can play a significant role in bringing some maturity to the field. Research fields advance through research papers. In research papers, however, only a limited perspective could be provided about the field, its application potential, and the techniques required and already developed in the field. A book gives such a chance. I liked the idea that there will be a book that will try to unify the field by bringing in disparate topics already available in several papers that are not easy to find and understand. I was supportive of this book project even before I had seen any material on it. The project was a brilliant and a bold idea by two active researchers. Now that I have it on my screen, it appears to be even a better idea.

Multimedia started gaining recognition in 1990s as a field. Processing, storage, communication, and capture and display technologies had advanced enough that researchers and technologists started building approaches to combine information in multiple types of signals such as audio, images, video, and text. Multimedia computing and communication techniques recognize correlated information in multiple sources as well as insufficiency of information in any individual source. By properly selecting sources to provide complementary information, such systems aspire, much like human perception system, to create a holistic picture of a situation using only partial information from separate sources.

Data mining is a direct outgrowth of progress in data storage and processing speeds. When it became possible to store large volume of data and run different statistical computations to explore all possible and even unlikely correlations among data, the field of data mining was born. Data mining allowed people to hypothesize relationships among data entities and explore support for those. This field has been put to applications in many diverse domains and keeps getting more applications. In fact many new fields are direct outgrowth of data mining and it is likely to become a powerful computational tool.



Preface

Approximately 17 million people in the USA (6% of the population) and 140 million people worldwide (this number is expected to rise to almost 300 million by the year 2025) suffer from *diabetes mellitus*. Currently, there are a few dozens of commercialised devices for detecting blood glucose levels [1]. However, most of them are invasive. The development of a noninvasive method would considerably improve the quality of life for diabetic patients, facilitate their compliance for glucose monitoring, and reduce complications and mortality associated with this disease. Noninvasive and continuous monitoring of glucose concentration in blood and tissues is one of the most challenging and exciting applications of optics in medicine. The major difficulty in development and clinical application of optical noninvasive blood glucose sensors is associated with very low signal produced by glucose molecules. This results in low sensitivity and specificity of glucose monitoring by optical methods and needs a lot of efforts to overcome this difficulty.

A wide range of optical technologies have been designed in attempts to develop robust noninvasive methods for glucose sensing. The methods include infrared absorption, near-infrared scattering, Raman, fluorescent, and thermal gradient spectroscopies, as well as polarimetric, polarization heterodyning, photonic crystal, optoacoustic, optothermal, and optical coherence tomography (OCT) techniques [1-31].

For example, the polarimetric quantification of glucose is based on the phenomenon of optical rotatory dispersion, whereby a chiral molecule in an aqueous solution rotates the plane of linearly polarized light passing through the solution. The angle of rotation depends linearly on the concentration of the chiral species, the pathlength through the sample, and the molecule specific rotation. However, polarization sensitive optical technique makes it difficult to measure *in vivo* glucose concentration in blood through the skin because of the strong light scattering which causes light depolarization. For this reason, the anterior chamber of the eye has been suggested as a sight well suited for polarimetric measurements, since scattering in the eye is generally very low compared to that in other tissues, and a high correlation exists between the glucose in the blood and in the aqueous humor. The high accuracy of anterior eye chamber measurements is also due to the low concentration of optically active aqueous proteins within the aqueous humor.

On the other hand, the concept of noninvasive blood glucose sensing using the scattering properties of blood and tissues as an alternative to spectral absorption and polarization methods for monitoring of physiological glucose

concentrations in diabetic patients has been under intensive discussion for the last decade. Many of the considered effects, such as changing of the size, refractive index, packing, and aggregation of RBC under glucose variation, are important for glucose monitoring in diabetic patients. Indeed, at physiological concentrations of glucose, ranging from 40 to 400 mg/dl, the role of some of the effects may be modified, and some other effects, such as glucose penetration inside the RBC and the followed hemoglobin glycation, may be important [30-32].

Noninvasive determination of glucose was attempted using light scattering of skin tissue components measured by a spatially-resolved diffuse reflectance or NIR frequency-domain reflectance techniques. Both approaches are based on change in glucose concentration, which affects the refractive index mismatch between the interstitial fluid and tissue fibers, and hence reduces scattering coefficient. A glucose clamp experiment showed that reduced scattering coefficient measured in the visible range qualitatively tracked changes in blood glucose concentration for the volunteer with diabetes studied.

List of Figures

1.1	List of figure caption goes here	9
1.2	Short figure caption	10
1.3	The bar charts depict the different risk contributions	11



List of Tables

1.1	Now we are engaged (a_g^a) (a_g^a) in a great civil war, testing whether that nation, or any nation so conceived.	5
1.2	Short Table Caption	6



Contributors

Michael Aftosmis

NASA Ames Research Center
Moffett Field, California

Pratul K. Agarwal

Oak Ridge National Laboratory
Oak Ridge, Tennessee

Sadaf R. Alam

Oak Ridge National Laboratory
Oak Ridge, Tennessee

Gabrielle Allen

Louisiana State University
Baton Rouge, Louisiana

Martin Sandve Alnæs

Simula Research Laboratory and
University of Oslo, Norway
Norway

Steven F. Ashby

Lawrence Livermore National
Laboratory
Livermore, California

David A. Bader

Georgia Institute of Technology
Atlanta, Georgia

Benjamin Bergen

Los Alamos National Laboratory
Los Alamos, New Mexico

Jonathan W. Berry

Sandia National Laboratories
Albuquerque, New Mexico

Martin Berzins

University of Utah
Salt Lake City, Utah

Abhinav Bhatele

University of Illinois
Urbana-Champaign, Illinois

Christian Bischof

RWTH Aachen University
Germany

Rupak Biswas

NASA Ames Research Center
Moffett Field, California

Eric Bohm

University of Illinois
Urbana-Champaign, Illinois

James Bordner

University of California, San Diego
San Diego, California

George Bosilca

University of Tennessee
Knoxville, Tennessee

Greg L. Bryan

Columbia University
New York, New York

Marian Bubak

AGH University of Science and
Technology

Kraków, Poland

Andrew Canning

Lawrence Berkeley National
Laboratory
Berkeley, California

Jonathan Carter

Lawrence Berkeley National
Laboratory
Berkeley, California

Zizhong Chen

Jacksonville State University
Jacksonville, Alabama

Joseph R. Crobak

Rutgers, The State University of
New Jersey
Piscataway, New Jersey

Roxana E. Diaconescu

Yahoo! Inc.
Burbank, California

Peter Diener

Louisiana State University
Baton Rouge, Louisiana

Jack J. Dongarra

University of Tennessee, Knoxville,
Oak Ridge National Laboratory,
and
University of Manchester

John B. Drake

Oak Ridge National Laboratory
Oak Ridge, Tennessee

Kelvin K. Droegemeier

University of Oklahoma
Norman, Oklahoma

Stéphane Ethier

Princeton University
Princeton, New Jersey

Christoph Freundl

Friedrich–Alexander–Universität
Erlangen, Germany

Karl Förlinger

University of Tennessee
Knoxville, Tennessee

Al Geist

Oak Ridge National Laboratory
Oak Ridge, Tennessee

Michael Gerndt

Technische Universität München
Munich, Germany

Tom Goodale

Louisiana State University
Baton Rouge, Louisiana

Tobias Gradl

Friedrich–Alexander–Universität
Erlangen, Germany

William D. Gropp

Argonne National Laboratory
Argonne, Illinois

Robert Harkness

University of California, San Diego
San Diego, California

Albert Hartono

Ohio State University
Columbus, Ohio

Thomas C. Henderson

University of Utah
Salt Lake City, Utah

Bruce A. Hendrickson

Sandia National Laboratories
Albuquerque, New Mexico

Alfons G. Hoekstra

University of Amsterdam
Amsterdam, The Netherlands

Philip W. Jones

Los Alamos National Laboratory
Los Alamos, New Mexico

Laxmikant Kalé

University of Illinois
Urbana-Champaign, Illinois

Shoaib Kamil

Lawrence Berkeley National
Laboratory
Berkeley, California

Cetin Kiris

NASA Ames Research Center
Moffett Field, California

Uwe Küster

University of Stuttgart
Stuttgart, Germany

Julien Langou

University of Colorado
Denver, Colorado

Hans Petter Langtangen

Simula Research Laboratory and
University of Oslo, Norway

Michael Lijewski

Lawrence Berkeley National
Laboratory
Berkeley, California

Anders Logg

Simula Research Laboratory and
University of Oslo, Norway

Justin Luitjens

University of Utah
Salt Lake City, Utah

Kamesh Madduri

Georgia Institute of Technology
Atlanta, Georgia

Kent-Andre Mardal

Simula Research Laboratory and

University of Oslo, Norway

Satoshi Matsuoka

Tokyo Institute of Technology
Tokyo, Japan

John M. May

Lawrence Livermore National
Laboratory
Livermore, California

Celso L. Mendes

University of Illinois
Urbana-Champaign, Illinois

Dieter an Mey

RWTH Aachen University
Germany

Tetsu Narumi

Keio University
Japan

Michael L. Norman

University of California, San Diego
San Diego, California

Boyana Norris

Argonne National Laboratory
Argonne, Illinois

Yousuke Ohno

Institute of Physical and Chemical
Research (RIKEN)
Kanagawa, Japan

Leonid Oliker

Lawrence Berkeley National
Laboratory
Berkeley, California

Brian O'Shea

Los Alamos National Laboratory
Los Alamos, New Mexico

Christian D. Ott

University of Arizona
Tucson, Arizona

James C. Phillips
University of Illinois
Urbana-Champaign, Illinois

Simon Portegies Zwart
University of Amsterdam,
Amsterdam, The Netherlands

Thomas Radke
Albert-Einstein-Institut
Golm, Germany

Michael Resch
University of Stuttgart
Stuttgart, Germany

Daniel Reynolds
University of California, San Diego
San Diego, California

Ulrich Rde
Friedrich–Alexander–Universitt
Erlangen, Germany

Samuel Sarholz
RWTH Aachen University
Germany

Erik Schnetter
Louisiana State University
Baton Rouge, Louisiana

Klaus Schulten
University of Illinois
Urbana-Champaign, Illinois

Edward Seidel
Louisiana State University
Baton Rouge, Louisiana

John Shalf
Lawrence Berkeley National
Laboratory
Berkeley, California

Bo-Wen Shen

NASA Goddard Space Flight Center
Greenbelt, Maryland

Ola Skavhaug
Simula Research Laboratory and
University of Oslo, Norway

Peter M.A. Sloom
University of Amsterdam
Amsterdam, The Netherlands

Erich Strohmaier
Lawrence Berkeley National
Laboratory
Berkeley, California

Makoto Taiji
Institute of Physical and Chemical
Research (RIKEN)
Kanagawa, Japan

Christian Terboven
RWTH Aachen University,
Germany

Mariana Vertenstein
National Center for Atmospheric
Research
Boulder, Colorado

Rick Wagner
University of California, San Diego
San Diego, California

Daniel Weber
University of Oklahoma
Norman, Oklahoma

James B. White, III
Oak Ridge National Laboratory
Oak Ridge, Tennessee

Terry Wilmarth
University of Illinois
Urbana-Champaign, Illinois

Symbols

Symbol Description

α	To solve the generator maintenance scheduling, in the past, several mathematical techniques have been applied.	$\theta\sqrt{abc}$	annealing and genetic algorithms have also been tested. This paper presents a survey of the literature
σ^2	These include integer programming, integer linear programming, dynamic programming, branch and bound etc.	ζ	over the past fifteen years in the generator
\sum	Several heuristic search algorithms have also been developed. In recent years expert systems,	∂	maintenance scheduling. The objective is to
abc	fuzzy approaches, simulated	sdf	present a clear picture of the available recent literature
		ewq	of the problem, the constraints and the other aspects of
		$bvcn$	the generator maintenance schedule.





Part I

**This is What a Part Would
Look Like**



1

Basic Concepts

CONTENTS

1.1	Introduction	3
1.1.1	A component part	4
1.2	Record Linkage Model	8
1.2.1	A component part	8
1.2.1.1	H3 A component part	8
1.3	Glossary	12

A component part for an electronic item is manufactured at one of three different factories, and then delivered to the main assembly line. Of the total number supplied, factory A supplies 50%, factory B 30%, and factory C 20%. Of the components manufactured at factory A, 1% are faulty and the corresponding proportions for factories B and C are 4% and 2% respectively. A component is picked at random from the assembly line. What is the probability that it is faulty?

1.1 Introduction

The term reliability usually refers to the probability that a component or system will operate satisfactorily either at any particular instant at which it is required or for a certain length of time. Fundamental to quantifying reliability is a knowledge of how to define, assess and combine probabilities [1]. This may hinge on identifying the form of the variability which is inherent in most processes. If all components had a fixed known lifetime there would be no need to model reliability.

1. A component part for an electronic item is manufactured at one of three different factories.
2. A component part x for an electronic item is manufactured at one of three different factories.

- (a) A component part for an electronic item is manufactured at one of three different factories.
 - (b) A component part x for an electronic item is manufactured at one of three different factories.
 - i. A component part for an electronic item is manufactured at one of three different factories.
 - ii. A component part x for an electronic item is manufactured at one of three different factories.
 - iii. A component part for an electronic item is manufactured at one of three different factories.
 - iv. A component part 1, 2, 3, 4 for an electronic item is manufactured at one of three different factories.
 - v. A component part for enumerate list of an electronic item is manufactured at one of three different factories.
 - (c) A component part for an electronic item is manufactured at one of three different factories.
 - (d) A component part 1, 2, 3, 4 for an electronic item is manufactured at one of three different factories.
 - (e) A component part for enumerate list of an electronic item is manufactured at one of three different factories.
- 3. A component part for an electronic item is manufactured at one of three different factories.
 - 4. A component part 1, 2, 3, 4 for an electronic item is manufactured at one of three different factories.
 - 5. A component part for enumerate list of an electronic item is manufactured at one of three different factories.

1.1.1 A component part

A component part for an electronic item is manufactured at one of three different factories, and then delivered to the main assembly line. Of the total number supplied, factory A supplies 50%, factory B 30%, and factory C 20%. Of the components manufactured at factory A, 1% are faulty and the corresponding proportions for factories B and C are 4% and 2% respectively. A component is picked at random from the assembly line. What is the probability that it is faulty [4]? A component part for an electronic item is manufactured at one of three different factories, and then delivered to the main assembly line. Of the total number supplied, factory A supplies 50%, factory B 30%, and factory C 20%. Of the components manufactured at factory A, 1% are faulty and the corresponding proportions for factories B and C are 4% and 2% respectively. A component is picked at random from the assembly line. What is the probability that it is faulty? A component part for an electronic item is manufactured at one of three different factories, and then delivered to the main

TABLE 1.1

Now we are engaged (a_g^a) (a_g^a) in a great civil war, testing whether that nation, or any nation so conceived.

Scene	Reg. fts.	Hor. fts.	Ver. fts.
Ball	19, 221	4, 598	3, 200
Pepsi ^a	46, 281	6, 898	5, 400
Keybrd ^b	27, 290	2, 968	3, 405
Pepsi	14, 796	9, 188	3, 209

assembly line. Of the total number supplied, factory A supplies 50%, factory B 30%, and factory C 20%. Of the components manufactured at factory A, 1% are faulty and the corresponding proportions for factories B and C are 4% and 2% respectively. A component is picked at random from the assembly line. What is the probability that it is faulty?

“A Process is a structured, measured set of activities designed to produce a specific output for a particular customer or market—
A process is thus a specific ordering of work activities across time and space, with a beginning, an end, and clearly defined inputs and outputs: a structure for action.”

Thomas Davenport
Senior Adjutant to the Junior Marketing VP

MultiRelational k -Anonymity. Most works on k -anonymity focus on anonymizing a single data table; however, a real-life [2] database usually contains multiple relational tables. This has proposed a privacy model called *MultiR k -anonymity* to ensure k -anonymity on multiple relational tables. Their model assumes that a relational database contains a person-specific table PT and a set of tables T_1, \dots, T_n , where PT contains a person identifier Pid and some sensitive attributes, and T_i , for $1 \leq i \leq n$, contains some foreign keys, some attributes in QID , and sensitive attributes. The general privacy notion is to ensure that for each record owner o contained in the join of all tables $PT \bowtie T_1 \bowtie \dots \bowtie T_n$, there exists at least $k - 1$ other record owners share the same QID with o . It is important to emphasize that the k -anonymization is applied at the *record owner* level, not at the *record* level in traditional k -anonymity. This idea is similar to (X, Y) -anonymity, where $X = QID$ and $Y = \{Pid\}$.

Most works on k -anonymity focus on anonymizing a single data table; however, a real-life [2] database usually contains multiple relational tables.

This has proposed a privacy model called *MultiR k-anonymity* to ensure k -anonymity on multiple relational tables. Their model assumes that a relational database contains a person-specific table PT and a set of tables T_1, \dots, T_n , where PT contains a person identifier Pid and some sensitive attributes, and T_i , for $1 \leq i \leq n$, contains some foreign keys, some attributes in QID , and sensitive attributes. The general privacy notion is to ensure that for each record owner o contained in the join of all tables $PT \bowtie T_1 \bowtie \dots \bowtie T_n$, there exists at least $k - 1$ other record owners share the same QID with o . It is important to emphasize that the k -anonymization is applied at the *record owner* level, not at the *record* level in traditional k -anonymity. This idea is similar to (X, Y) -anonymity, where $X = QID$ and $Y = \{Pid\}$.

A component part for an electronic item is manufactured at one of three different factories, and then delivered to the main assembly line. Of the total number supplied, factory A supplies 50%, factory B 30%, and factory C 20%. Of the components manufactured at factory A, 1% are faulty and the corresponding proportions for factories B and C are 4% and 2% respectively. A component is picked at random from the assembly line. What is the probability that it is faulty?

In most literature on PPDP, they [6] consider a more relaxed, yet more practical, notion of privacy protection by assuming limited attacker's background knowledge. Below, the term "victim" refers to the record owner being linked. We can broadly classify linking models to two families.

A component part for an electronic item is [3] manufactured at one of three different factories, and then delivered to the main assembly line. Of the total number supplied, factory A supplies 50%, factory B 30%, and factory C 20%. Of the components manufactured at factory

TABLE 1.2

Now we are engaged (a_g^a) (a_g^a) in a great civil war, testing whether that nation, or any nation so conceived.

Scene	Reg. fts.	Hor. fts.	Ver. fts.
Table Head			
Ball	19, 221	4, 598	3, 200
Pepsi	46, 281	6, 898	5, 400
Keybrd	27, 290	2, 968	3, 405
Pepsi	14, 796	9, 188	3, 209

A, 1% are faulty and the corresponding proportions for factories B and C are 4% and 2% respectively.

One family considers a privacy threat occurs when an attacker is able to link a record owner to a record in a published data table, to a sensitive attribute in a published data table, or to the published data table itself. We call them *record linkage*, *attribute linkage*, and *table linkage*, respectively. In all types of linkages, we assume that the attacker knows the *QID* of the victim. In record and attribute linkages, we further assume that the attacker knows the presence of the victim's record in the released table, and seeks to identify the victim's record and/or sensitive information from the table [10]. In table linkage, the attack seeks to determine the present or absent of the victim's record in the released table. A data table is considered to privacy preserved if the table can effectively prevent the attacker from successfully performing these types of linkages on the table [7]. Sections 1.1-1.2 study this family of privacy models.

$$\text{var} \hat{\Delta} = \sum_{j=1}^t \sum_{k=j+1}^t \text{var} (\hat{\alpha}_j - \hat{\alpha}_k) = \sum_{j=1}^t \sum_{k=j+1}^t \sigma^2 (1/n_j + 1/n_k). \quad (1.1)$$

An obvious measure of imbalance is just the difference in the number of times the two treatments are allocated

$$D_n = \mathcal{M}|n_A - n_B|. \quad (1.2)$$

For rules such as deterministic allocation, for which the expected value of this difference can be calculated, we obtain the population value \mathcal{D}_n .

Box Title Here

Another family aims at achieving the *uninformative principle*: The published table should provide the attacker with little additional information beyond the background knowledge. There should not be a large difference between the prior and posterior beliefs; otherwise, there is a privacy threat [5, 6]. Many privacy models in this family are designed for statistical database and do not distinguish attributes in T into *QID*, but some of them could also thwart record, attribute, and table linkages. Section 1.1 studies this family of privacy models.

Let m be a prime number. With the addition and multiplication as defined above, Z_m is a field.

Theorem 1 *Let m be a prime number. With the addition and multiplication as defined above, Z_m is a field.*

Proof 1 *Most of the proof of this theorem is routine. It is clear that $0 \in Z_m$ and $1 \in Z_m$ are the zero element and identity element. If $a \in Z_m$ and $a \neq 0$,*

then $m - a$ is the additive inverse of a . If $a \in Z_m$ and $a \neq 0$, then the greatest common divisor of a and m is 1, and hence there exist integers s and t such that $sa + tm = 1$. Thus $sa = 1 - tm$ is congruent to 1 modulo m . Let s^* be the integer in Z_m congruent to s modulo m . Then we also have $s^*a \equiv 1 \pmod{m}$. Hence s^* is the multiplicative inverse of a modulo m . Verification of the rest of the field properties is now routine.

1.2 Record Linkage Model

In the privacy attack of *record linkage*, some value qid on QID identifies a small number of records in the released table T , called a *group*. If the victim's QID matches the value qid , the victim is vulnerable to being linked to the small number of records in the group [8]. In this case, the attacker faces only a small number of possibilities for the victim's record, and with the help of additional knowledge, there is a chance that the attacker could uniquely identify the victim's record from the group.

1.2.1 A component part

A component part for an electronic item is manufactured at one of three different factories, and then delivered to the main assembly line. Of the total number supplied, factory A supplies 50%, factory B 30%, and factory C 20%. Of the components manufactured at factory A, 1% are faulty and the corresponding proportions for factories B and C are 4% and 2% respectively. A component is picked at random from the assembly line. What is the probability that it is faulty?

1.2.1.1 H3 A component part

A component part for an electronic item is manufactured at one of three [9] different factories, and then delivered to the main assembly line. Of the total number supplied, factory A supplies 50%, factory B 30%, and factory C 20%. Of the components manufactured at factory A, 1% are faulty and the corresponding proportions for factories B and C are 4% and 2% respectively. A component is picked at random from the assembly line. What is the probability that it is faulty?

A fundamental notion [10] is that of a subspace of F^n . Let V be a nonempty subset of F^n . Then V is a *subspace* of F^n provided V is closed under vector addition and scalar multiplication, that is,

- (a) For all u and v in V , $u + v$ is also in V .
- (b) For all u in V and c in F , cu is in V .


**FIGURE 1.1**

Figure caption goes here.

Let u be in the subspace V . Because $0u = 0$, it follows that the zero vector is in V . Similarly, $-u$ is in V for all u in V . A simple example of a subspace of F^n is the set of all vectors $(0, a_2, \dots, a_n)$ with first coordinate equal to 0. The zero vector itself is a subspace.

Definition 1 Let $u^{(1)}, u^{(2)}, \dots, u^{(m)}$ be vectors in F^n , and let c_1, c_2, \dots, c_m be scalars. Then the vector

$$c_1 u^{(1)} + c_2 u^{(2)} + \dots + c_m u^{(m)}$$

is called a *linear combination* of $u^{(1)}, u^{(2)}, \dots, u^{(m)}$. If V is a subspace of F^n , then V is closed under vector addition and scalar multiplication, and it follows easily by induction that a linear combination of vectors in V is also a vector in V . Thus *subspaces are closed under linear combinations*; in fact, this can be taken as the defining property of subspaces. The vectors $u^{(1)}, u^{(2)}, \dots, u^{(m)}$ *span* V (equivalently, form a *spanning set* of V) provided every vector in V is a linear combination of $u^{(1)}, u^{(2)}, \dots, u^{(m)}$. The zero vector can be written as a linear combination of $u^{(1)}, u^{(2)}, \dots, u^{(m)}$ with all scalars equal to 0; this is a *trivial linear combination*. The vectors $u^{(1)}, u^{(2)}, \dots, u^{(m)}$ are *linearly dependent* provided there are scalars c_1, c_2, \dots, c_m , not all of which are zero, such that

$$c_1 u^{(1)} + c_2 u^{(2)} + \dots + c_m u^{(m)} = 0,$$

that is, the zero vector can be written as a *nontrivial linear combination* of


**FIGURE 1.2**

Figure caption goes here. Figure caption goes here. Figure caption goes here.
 Figure caption goes here. Figure caption goes here. Figure caption goes here.

$u^{(1)}, u^{(2)}, \dots, u^{(m)}$. For example, the vectors $(1, 4)$, $(3, -1)$, and $(3, 5)$ in \mathfrak{R}^2 are linearly dependent since

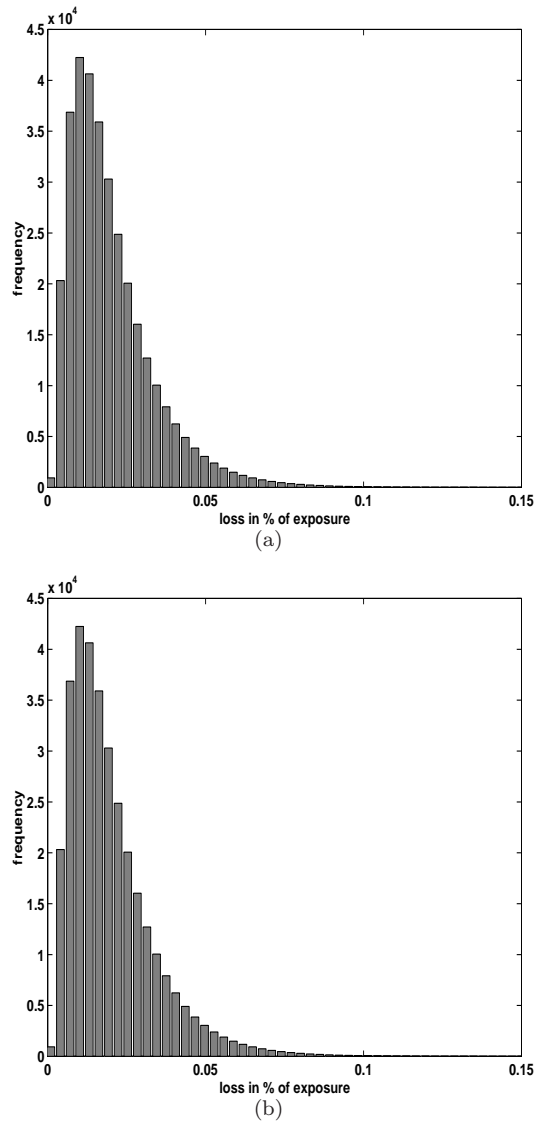
$$3(1, 4) + 1(3, -2) - 2(3, 5) = (0, 0).$$

Vectors are *linearly independent* provided they are not linearly dependent. The vectors $u^{(1)}, u^{(2)}, \dots, u^{(m)}$ are a *basis* of V provided they are linearly independent and span V . By an *ordered basis* we mean a basis in which the vectors of the basis are listed in a specified order; to indicate that we have an ordered basis we write $(u^{(1)}, u^{(2)}, \dots, u^{(m)})$. A spanning set S of V is a *minimal spanning set* of V provided that each set of vectors obtained from S by removing a vector is not a spanning set for V . A linearly independent set S of vectors of V is a *maximal linearly independent set of vectors of V* provided that for each vector w of V that is not in S , $S \cup \{w\}$ is linearly dependent (when this happens, w must be a linear combination of the vectors in S). \square

In addition to matrix addition, subtraction, and multiplication, there is one additional operation that we define now. It's perhaps the simplest of them all. Let $A = [a_{ij}]$ be an m by n matrix and let c be a number [3]. Then the matrix $c \cdot A$, or simply cA , is the m by n matrix obtained by multiplying each entry of A by c :

$$cA = [ca_{ij}].$$

The matrix cA is called a *scalar multiple* of A .

**FIGURE 1.3**

The bar charts depict the different risk contributions (top: 99% quantile, bottom: 99.9% quantile) of the business areas of a bank. The black bars are based on a Var/Covar approach, the white ones correspond to shortfall risk.

Think About It...

Commonly thought of as the first modern computer, ENTAC was built in 1944. It took up more space than an 18-wheeler's tractor trailer and weighed more than 17 Chevrolet Camaros. It consumed 140,000 watts of electricity while executing up to 5,000 basic arithmetic operations per second. One of today's popular microprocessors, the 486, is built on a tiny piece of silicon about the size of a dime.

With the continual expansion of capabilities, computing power will eventually exceed the capacity for human comprehension or human control.

The Information Revolution
Business Week

1.3 Glossary

360 Degree Review: Performance review that includes feedback from superiors, peers, subordinates, and clients.

Abnormal Variation: Changes in process performance that cannot be accounted for by typical day-to-day variation. Also referred to as non-random variation.

Acceptable Quality Level (AQL): The minimum number of parts that must comply with quality standards, usually stated as a percentage.

Activity: The tasks performed to change inputs into outputs.

Adaptable: An adaptable process is designed to maintain effectiveness and efficiency as requirements change. The process is deemed adaptable when there is agreement among suppliers, owners, and customers that the process will meet requirements throughout the strategic period.

Bibliography

- [1] G. Bontempi and Y. Le Borgne. An adaptive modular approach to the mining of sensor network data. In *Proceedings of the Workshop on Data Mining in Sensor Networks, SIAM SDM*, pages 3–9. SIAM Press, 2005.
- [2] KI Diamantaras and SY Kung. *Principal component neural networks: theory and applications*. John Wiley & Sons, Inc. New York, NY, USA, 1996.
- [3] A. Hyvarinen, J. Karhunen, and E. Oja. *Independent Component Analysis*. J. Wiley New York, 2001.
- [4] M. Ilyas, I. Mahgoub, and L. Kelly. *Handbook of Sensor Networks: Compact Wireless and Wired Sensing Systems*. CRC Press, Inc. Boca Raton, FL, USA, 2004.
- [5] A. Jain and E.Y. Chang. Adaptive sampling for sensor networks. *ACM International Conference Proceeding Series*, pages 10–16, 2004.
- [6] I.T. Jolliffe. *Principal Component Analysis*. Springer, 2002.
- [7] S. Madden, M.J. Franklin, J.M. Hellerstein, and W. Hong. TAG: a Tiny AGgregation Service for Ad-Hoc Sensor Networks. In *Proceedings of the 5th ACM Symposium on Operating System Design and Implementation (OSDI)*, volume 36, pages 131 – 146. ACM Press, 2002.
- [8] S.R. Madden, M.J. Franklin, J.M. Hellerstein, and W. Hong. TinyDB: an acquisitional query processing system for sensor networks. *ACM Transactions on Database Systems (TODS)*, 30(1):122–173, 2005.
- [9] K.V. Mardia, J.T. Kent, J.M. Bibby, et al. *Multivariate analysis*. Academic Press New York, 1979.
- [10] Y. Yao and J. Gehrke. The cougar approach to in-network query processing in sensor networks. *ACM SIGMOD Record*, 31(3):9–18, 2002.



Index

basis, 10
 ordered, 10

linear combination, 9
 nontrivial, 9
 trivial, 9

linearly independent, 10
 maximal, 10

matrix
 scalar multiple of, 10
 scalar multiplication, 10

spanning set, 9
 minimal, 10

subspace, 8

vector space
 subspace of, 8